

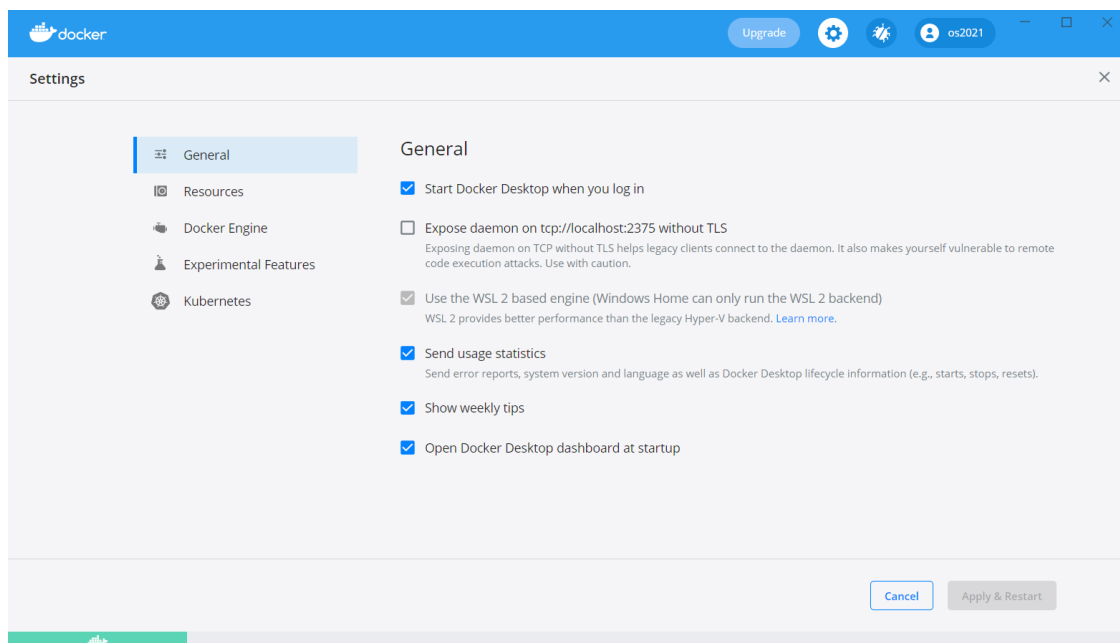
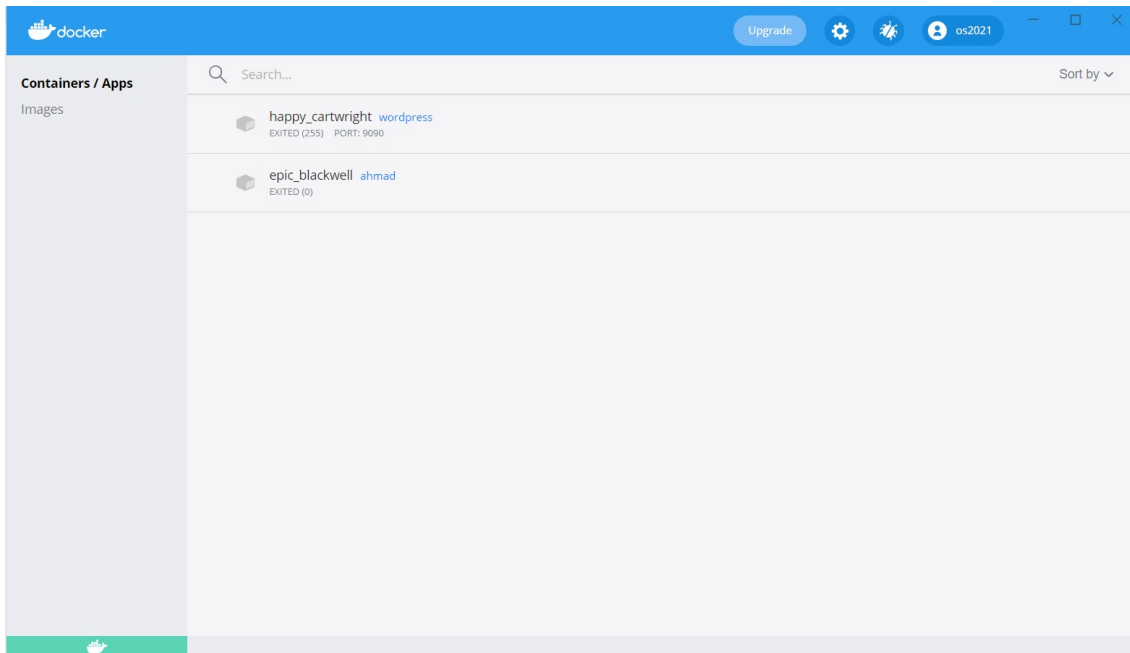
Before you start, make sure the following are installed on your laptop

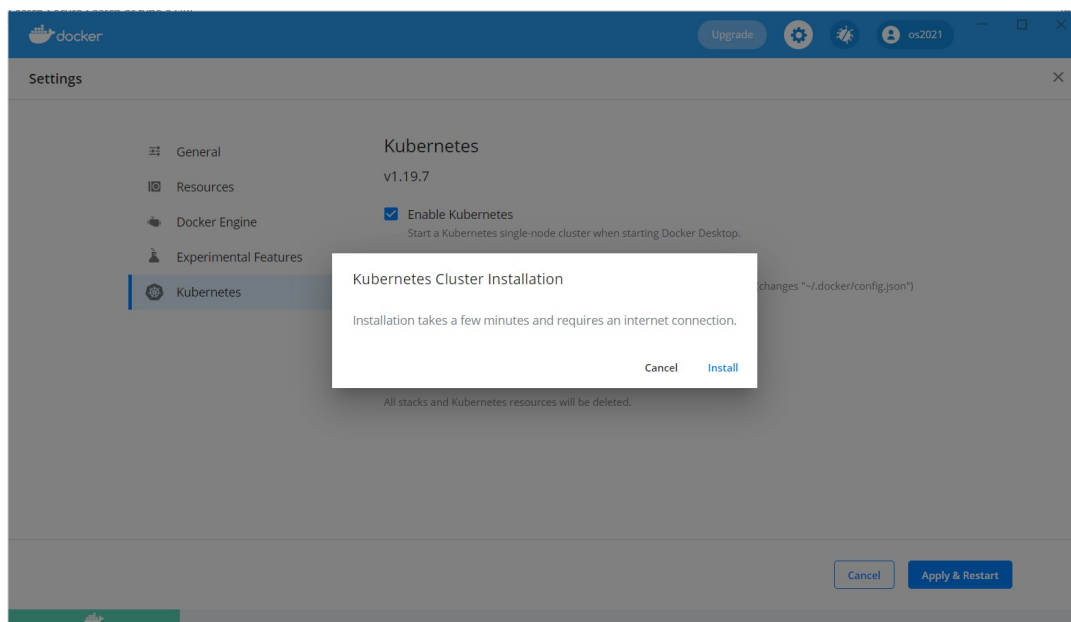
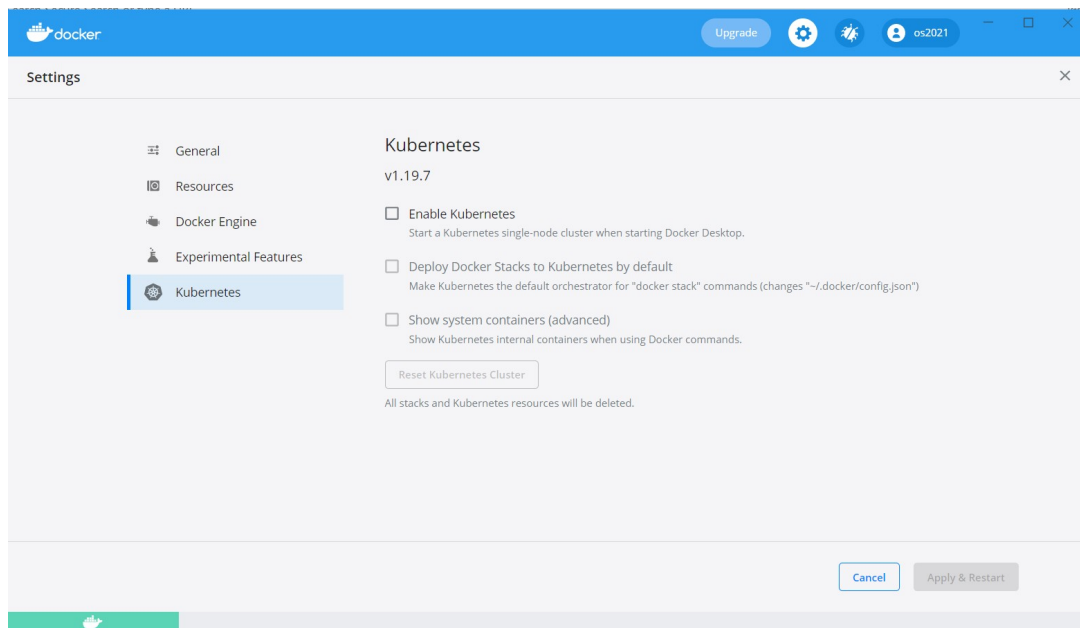
1. Kubernetes ( **Minikube** - <https://www.youtube.com/watch?v=8h4FoWK7tIA> )

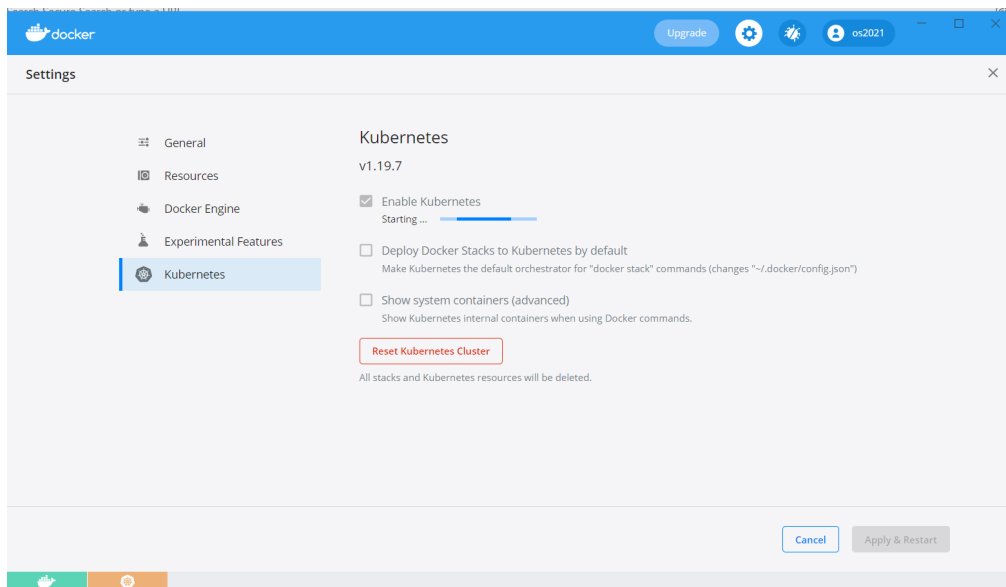
Can be installed easily if you have docker installed already on your system

2. Kube controller ( **kubectrl** - <https://minikube.sigs.k8s.io/docs/start/> )

On **Windows**, just follow the following pictures







1. To see a list of nodes ( currently we only have only one node since we used **minikube** )

***\$ kubectl get nodes***

```
ahmad@Batanouni:~$ kubectl get nodes
NAME        STATUS    ROLES    AGE   VERSION
minikube    Ready     master   6d23h v1.19.4
```

if we installed **kubeadm** instead, on **multiple** nodes we would get something like this

```
root@kubemaster:/home/osboxes# kubectl get nodes
NAME           STATUS    ROLES    AGE   VERSION
kubemaster     Ready     master   7m    v1.9.4
kubenode1      Ready     <none>    4m    v1.9.4
kubenode2      Ready     <none>    4m    v1.9.4
```

2. To list all the **Pods**

***\$ kubectl get pods***

3. To create a new pod ( pod is the smallest unit in Kubernetes which holds inside a docker container )

***\$ kubectl run nginx --image=nginx***

4. To view more information about the pod ... like showing on which node it was deployed, run

***\$ kubectl describe pods***

By: [Ahmad ElBatanouni](#)

5. To delete a pod

**\$ kubectl delete pod *podName***

6. create a pod from a **yaml** file ( files are needed for kubernetes to automatically restart failing pods )

**\$ kubectl create -f pod-definition.yml**

**All YAML files in kubernetes follow the following structure**

```
apiVersion:
kind:
metadata:

spec:
```

1. apiVersion

Can be one of [ "v1", "apps/v1" ] depending on what object we're trying to create

Pod	v1
Deployment	apps/v1
Service	v1

2. kind

Can be one of [ "Pod", "Service", "Deployment" ]

3. metadata

add information about the item we are trying to create

4. spec

in case of creating a

4.1 **POD**

you add information about the container/s you want to include in the pod and their resources ( limits )

4.2 **Deployment**

you specify the number of replicas (pods) that must be running all of the time, and also specify a template

Samples for **POD** & **Deployment** files are in the samples folder.

# Pods

Let's start by creating our first POD, go and fetch the template from the top of this page

```
apiVersion:  
kind:  
metadata:  
  
spec:
```

Since we're trying to create a POD, the suitable **apiVersion** is **v1** ... we know that from the table in the page above

```
apiVersion: v1  
kind:  
metadata:  
  
spec:
```

Next, we set the **kind** to **pod**

```
apiVersion: v1  
kind: pod  
metadata:  
  
spec:
```

Next, we set the **metadata** .. here we can specify 2 attributes ( **name** & **label** )

- **name** .. is just the name of the entity we're trying to create ( the POD in this case )
- **label** .. allows us to give the created POD/s a label that we can use later to perform some kind of operation on **all the PODs having this label**

```
apiVersion: v1  
kind: pod  
metadata:  
  name: any-name-that-you-like  
  labels:  
    label-name-1: label-value  
    another-label: my-app-name  
spec:
```

Note that we can add more than one label.

Finally, we move on to **spec** where we describe the kind of container/s we want to include in the pod

we do that by adding a child called **containers**

```
apiVersion: v1
kind: pod
metadata:
  name: any-name-that-you-like
  labels:
    label-name-1: label-value
    another-label: my-app-name
spec:
```

**containers:**

this attribute is in fact an **array** where you can add details of **more than one** container. Every item in this array is preceded with a dash ( - ) ... for example

```
apiVersion: v1
kind: pod
metadata:
  name: any-name-that-you-like
  labels:
    label-name-1: label-value
    another-label: my-app-name
spec:
  containers:
    - details of the first container
    - details of the second container
    - details of the third container
```

Note that the above syntax is invalid ... it is just used to explain how arrays in Yaml are created.

Next we specify the details of every item ( container ) like the **image name**, **container name**, **resource limits** ( like how much ram and CPU can it use )

```
apiVersion: v1
kind: pod
metadata:
  name: any-name-that-you-like
  labels:
    label-name-1: label-value
    another-label: my-app-name
spec:
  containers:
    - image: nginx
      name: nginx-container
      resources:
        limits:
          memory: "512Mi"
          cpu: "1"
```

Now that's we're done of the pod file ... we can create the pod itself using the following command

**\$ kubectl create -f filename.yml**

---

To make a pod accessible to the outside world, and not just withing the Kubernetes cluster, you'd use the **port-forward** command, giving a name of a deployment along with a mapping between a port on the host and a port inside the pod

**\$ kubectl port-forward [name-of-pod] 8080:80**

# Deployments

Another concept/entity is the **deployment controller** is a way to **manage** pods, by introducing the following capabilities

- **Deploy/Create** multiple instances of your application
- **Self-healing** making sure the **desired number** of pods is always there, if a pod goes down, the deployment controller spins up a new one
- **Update** the pods/containers with newer images using what's called "rolling update" which doesn't stop all the pods immediately and pull the new image/s ... **instead** it stops the existing pods either one by one or at a controlled rate resulting what's called a zero-downtime update.
- **Go back/Rollback** to an older deployment.
- **Scale up/down** the number of pods you have in the cluster.

To create a new deployment controller, we use the same template as before, in addition to that we **specify** 2 attributes under **spec** ... the first is **replicas** and the second is **template**

**replicas** specify the number of pods we need ... the deployment then makes sure the specified number of pods are always up and running ... and if for some reason one of the pods fails ... it spins up a new pod ( with the characteristics which we describe inside of the second attribute **template** ) ... **it also deletes pods if needed**.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-first-deployment
```

```
spec:
  replicas: 3
  template:
```

**Details of the pod to  
create**

Again the above syntax ( **details of the pod to create** ) is not valid ... it only shows that we use the value here to spin up new PODs when needed.



Those details are exactly the same as the details we specified in the pod creation file .. except we ignore the first 2 attributes ( apiVersion & kind )

So we can simply open the first pod file and **copy all what's under ( metadata & spec )** and paste them here in the box

*pod.yaml*

```
apiVersion: v1
kind: pod
metadata:
  name: any-name-that-you-like
  labels:
    label-name-1: label-value
    another-label: my-app-name
spec:
  containers:
    - image: nginx
      name: nginx-container
      resources:
        limits:
          memory: "512Mi"
          cpu: "1"
```

*deployment.yaml*

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment-controller
spec:
  replicas: 3
  template:
    metadata:
      name: any-name-that-you-like
      labels:
        label-name-1: label-value
        another-label: my-app-name
    spec:
      containers:
        - image: nginx
          name: nginx-container
          resources:
            limits:
              memory: "512Mi"
              cpu: "1"
```

As usual, create the deployment using

**\$ kubectl create -f filename.yml**

By: [Ahmad ElBatanouni](#)

To see the deployment controller we've just created

***\$ kubectl get deployment***

The deployment in turn creates a replica set with the same name, check that by running

***\$ kubectl get rs***

Another important attribute that we can specify under **spec** is called **selector** which is used by the deployment controller to track all the pods matching that label

***deployment.yml***

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment

spec:
  selector:
    matchLabels:
      my-label: some-value
  replicas: 3
  template:
    metadata:
      name: any-name-that-you-like
      labels:
        my-label: some-value
    spec:
      containers:
        - image: nginx
          name: nginx-container
```

To **update** the existing pods with a different version of the container image, we have 2 strategies

**1. Recreate:** Stops all the pods at once, and recreates them using the new image .. the **problem** here is that users will be unable to access the application during the update time.

**2. RollingUpdate** ( Default ): stops and updates the existing pods either one by one or at a controlled rate so that users can access the application on other pods.

We can then update the yaml file with one of those 2 strategies **along with the update we want to do to the container image**

By: [Ahmad ElBatanouni](#)

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment

spec:
  selector:
    matchLabels:
      my-label: some-value
  replicas: 3
  strategy:
    type: Recreate
  template:

```

```

  metadata:
    name: any-name-that-you-like
    labels:
      my-label: some-value
  spec:
    containers:
      - image: nginx:1.16.1
      name: nginx-container

```

Apply the changes using

**\$ kubectl **apply** -f deployment.yml**

Now try to change the number of replicas ( either increase or decrease them ) and again run:

**\$ kubectl **apply** -f deployment.yml**

Let's assume that updating to the new image wasn't the correct decision, something is just wrong with that image ... we want to move back ( **rollback** ) to the old image ( in Kubernetes world it's called the old **revision** )

**\$ kubectl **rollback** **undo** deployment/my-deployment**

If you now run **\$kubectl get pods** you can see the pods are brought down one by one and new pods are created instead.

To **scale up/down** the number of replicas in a certain deployment, you'd simply do something like this

**\$ kubectl **scale** deployment [name-of-depl] --replicas=5**

Finally, to make a pod/deployment accessible to the outside world, and not just within the Kubernetes cluster, you'd use the **port-forward** command, giving a name of a deployment along with a mapping between a port on the host and a port on

**\$ kubectl port-forward [name-of-pod] 8080:80**

**\$ kubectl port-forward [name-of-deployment] 8080:80**

**8080 is the host port**  
**80 is the pod/deployment port**

### **Assignment**

Create a Kubernetes deployment file that creates 2 replicas of **wordpress**

1. Show the names of all the replicas
2. Show an example of how you can access wordpress from the browser ( P.S. port forwarding ).
3. Increase the number of replicas to 3 then show again the names of all of them.
4. Change the image from **wordpress** to **nginx**